



FIT2100 Laboratory #2  
Managing Larger C Programs  
&  
Running I/O System Calls in C  
Week 4 Semester 2 2021

August 9, 2021

**Revision Status:**

Written by Dr. Jojo Wong,  
Updated by Dr. Charith Jayasekara, Aug 2021.

## Acknowledgements

The content presented in Sections 2.1 and 2.2 in this laboratory were adapted from the courseware of FIT3042 prepared by Robyn McNamara.

The content presented in Section 2.3 and the laboratory tasks (Section 3) were adapted from David Curry's texts.

- David A. Curry (1989). *C on the UNIX System*, O'Reilly.
- David A. Curry (1996). *UNIX Systems Programming for SVR4*, O'Reilly.

## Contents

<b>1</b>	<b>Background</b>	<b>4</b>
<b>2</b>	<b>Pre-class Preparation</b>	<b>4</b>
2.1	Managing C programs . . . . .	4
2.1.1	Organising your source files . . . . .	4
2.1.2	Building complex C programs . . . . .	5
2.1.3	The <code>makefile</code> format . . . . .	6
2.1.4	Running the <code>make</code> command . . . . .	7
2.1.5	Defining macros in <code>makefile</code> . . . . .	8
2.2	I/O System Calls in C . . . . .	10
2.2.1	File descriptors . . . . .	10
2.2.2	Opening and creating files . . . . .	11
2.2.3	Reading and writing files . . . . .	12
2.2.4	Using <code>write()</code> with strings . . . . .	13

2.2.5	Moving around files . . . . .	14
2.3	Pre-class Quiz Checkpoint (2 marks) . . . . .	15
<b>3</b>	<b>Laboratory Tasks (8 marks)</b>	<b>16</b>
3.1	Task 1 (2 marks) . . . . .	16
3.2	Task 2 (3 marks) . . . . .	17
3.3	Task 3 (3 marks) . . . . .	17
3.4	Task 4 (Extension: 0 marks) . . . . .	18
<b>4</b>	<b>Postscript: Redirects and Pipes</b>	<b>18</b>

# 1 Background

In this second laboratory, there are two objectives. Firstly, you will learn how to manage complex C programs. Secondly, you will learn how to perform low-level I/O systems calls in C.

Before you attempt any of the tasks in this laboratory, create a folder named LAB02 under the FIT2100 folder (`~/Documents/FIT2100`). Save all your source files under this LAB02 folder.

Before attending the lab, you should:

- Complete your pre-class readings and preparations (Section 2)
- Complete your pre-class quiz
- Attempt the lab tasks (Section 3)

## 2 Pre-class Preparation

### 2.1 Managing C programs

Often times, working with a small programming project, you could probably get away with putting all your code in one source file. For much larger projects, you would need to organise your functions together into various source files by grouping related functions together.

A general *rule of thumb* is that your functions should not be much larger than a terminal screenful (approximately 25 lines of code), excluding header comments. If your functions become larger than this, you should consider breaking them down into smaller functions. (Your function names should be descriptive to help other coders to navigate your source files.)

#### 2.1.1 Organising your source files

One characteristic of the C Programming Language is that it encourages the idea of 'separate compilation.' Instead of putting all the code in one file and compiling that one file, C allows you to write many `.c` files and compile them separately when necessary. The usual way to group related C source files is to put them under the same directory.

**What goes into the .c file?** Anything that causes the compiler to generate code should go into the .c source file. A .c source file usually consists of:

- implementation of functions (i.e. function bodies)
- declaration of global variables — in order to set aside memory for global variables

Anything that is intended as a message to the compiler should go into the .h header file, which include:

- function prototypes
- struct definitions
- typedef statements
- #define statements

**Note:** Only .h files should ever be included using the #include statement — you should never include .c files.

For example, if your program has a .c file named `arithmetic.c`, you would place the function prototypes and other related statements in a header file named `arithmetic.h`, and #include "arithmetic.h" both at the top of `arithmetic.c` and at the top of any other file that needs access to those function prototypes. Including just the header information in each file rather than the entire functions themselves avoids a problem where the C compiler 'sees' the same function redefined multiple times.

### 2.1.2 Building complex C programs

A C program can potentially consist of hundreds of source files (both .c and .h files), and possibly spanning across many directories.

The standard command given below is used to compile a complex C program with multiple source files. However, *this approach does not scale!*

```
1 $ gcc sourcefile1.c sourcefile2.c ... -o myprogram
```

Whenever the program is rebuilt, this command recompiles every source file, which can be time consuming. It should only be recompiling those source files that have been modified instead of every single file. The solution to this is by utilising the `make` command in C.

`make` is a utility tool in C that automates the building of program binaries from source files. A number of built-in rules are defined, such that `make` knows how to compile source code (`.c` files) into object code (`.o` files). `make` requires a special file called `makefile` that specifies the source files and the targets (which include the executable code) to be build.

### 2.1.3 The makefile format

A `makefile` typically consists of one or more entries. Each entry consists of the following:

- a target (which usually is a file but not always);
- a list of dependencies (files which the target depends on);
- and the commands to execute based on the target and its dependencies.

The basic format for each entry in a `makefile` looks like below:

```
1 <target>: [ <dependency> ]*  
2 [ <TAB> <command> <ENDLINE> ]+
```

(Note: There must be a `<TAB>` character at the start of each command defined under the target. Indentation using spaces will *not* work in a `makefile`.)

An example of a `makefile` <sup>1</sup>:

```
1 sourcefile1.o: sourcefile1.c headerfile1.h  
2 gcc -Wall -c sourcefile1.c
```

In the example above, the `-c` option is needed to create the corresponding object file (`.o`) for the given `.c` file. The `-Wall` option (stands for *Warnings: all*), which you might not have seen before, is a way to turn on extra warnings in `gcc`, which can be very useful for finding hidden problems in your code.

---

<sup>1</sup>The second line starts with a `<TAB>`, and not spaces.

**How and when is a target constructed?** Each entry in the `makefile` defines *how* and *when* to construct a target based on its dependencies. The dependencies are used by the `makefile` to determine when the target needs to be reconstructed (by re-compiling the source code).

Each file has a *timestamp* that indicates when the file was last modified.

The `make` command first checks the timestamp of the file (i.e. the target) and then the timestamp of its dependencies (which are also files). If the dependent files have a more recent timestamp than the target file, the command lines defined under the entry of that target are executed to update the target.

**Note:** The updates are done recursively. If any of the dependencies themselves are also targets (i.e. there is a separate entry in the `makefile` for each of the dependencies), the `make` command will check whether these dependent files need to be updated, before updating the initial target.

**How to create the executable code using targets?** The purpose of having the `makefile` is to create the executable code (i.e. `a.out`). The executable file is often the first target.

```
1 myprogram: sourcefile1.o sourcefile2.o sourcefile3.o
2 gcc -Wall sourcefile1.o sourcefile2.o sourcefile3.o -o myprogram
```

#### 2.1.4 Running the `make` command

Once you have a `makefile` written, you can make use of the `make` command to update your files whenever you make changes. There are two ways to run `make` as shown below.

```
1 $ make
2 $ make -f <makefilename>
```

If you just type `make` on the command prompt, the file called `makefile` in the current directory will be interpreted and the commands of the first target will be executed, provided that the first target has dependencies listed.

Alternatively, if you have more than one `makefile` with different file names, you will run the `make` command with the `-f` option by supplying the file name of the specific `makefile` that you would like to run.

### 2.1.5 Defining macros in makefile

You can use *macros* in your makefile, which allow you to define variables that can be substituted.

You can use the “=” operator to set a value to the macro, and use the “\$” operator when using the value of the macro<sup>2</sup>. (Note that the parentheses are required if the macro name has more than one character.)

**Important:** Don't get confused with the shell command prompt which can also be represented by a “\$” symbol.

```
1 CC = gcc
2 CFLAGS = -Wall -c
3
4 sourcefile1.o: sourcefile1.c headerfile1.h
5 $(CC) $(CFLAGS) sourcefile1.c
```

The example above demonstrates two common macros CC and CFLAGS. Below are some of the other common examples of macros.

Macro	Description
CC	the name of the compiler (gcc)
DEBUG	the debugging flag (-g)
CFLAGS	the flags used for compiling (e.g. -c, -Wall)
LFLAGS	the flags used for linking (e.g. -L, -I)
\$@	the file name of the target
\$^	the list of dependencies
\$<	the first item in the dependencies list

---

<sup>2</sup>If you have used shell variables in Bash before, makefile macros have a similar syntax.



**Putting it all together** Here is a sample makefile for creating the executable program called myprogram.

```
1 OBJS = sourcefile1.o sourcefile2.o sourcefile3.o
2 CC = gcc
3 CFLAGS = -Wall -c
4 LFLAGS = -Wall
5
6 myprogram: $(OBJS)
7     $(CC) $(LFLAGS) $(OBJS) -o myprogram
8
9 sourcefile1.o: sourcefile1.c headerfile1.h
10    $(CC) $(CFLAGS) sourcefile1.c
11
12 sourcefile2.o: sourcefile2.c headerfile2.h
13    $(CC) $(CFLAGS) sourcefile2.c
14
15 sourcefile3.o: sourcefile3.c headerfile3.h
16    $(CC) $(CFLAGS) sourcefile3.c
```

**Try creating a makefile yourself:**

Refer to Task 3 that you have completed for the Week 1 Tutorial. First, organise your code into three separate files: main.c, arithmetic.c, and arithmetic.h.

- main.c: consists of the main() function
- arithmetic.c: consists of the implementation of the four arithmetic functions (i.e. function bodies)
- arithmetic.h: consists of the function prototypes of the four arithmetic functions

Then, create a makefile and try to compile your program by using the make command.

## 2.2 I/O System Calls in C

*Note: You should complete the reading of this section before attempting the lab tasks in the next section.*

The Standard I/O Library (`stdio`) provides a collection of *high-level* routines to perform input and output (I/O). This enables C programmers to perform reading and writing data easily and efficiently.

In general, the `stdio` routines perform three important functions for programmers:

- Buffering on input and output data is performed automatically;
- Input and output conversions are performed using routines like `printf` and `scanf`;
- Input and output are also automatically formatted.

However, these functions such as buffering and I/O conversions are not always applicable. In the event where performing I/O directly from a device such as a disk drive, programmers would need to be able to determine the actual buffer size to be used instead of relying on the `stdio` routines.

Thus, a direct interface to the operating system is desirable and can be achieved by issuing *system calls* or using the *low-level* I/O interface. System calls are low-level OS/kernel functions.

### 2.2.1 File descriptors

When dealing with the low-level I/O interface, each of low-level I/O routines requires a valid “file descriptor” to be passed to them — which is used to reference a file to an open stream for I/O.

A file descriptor is simply a small integer, which can be allocated from a file index table maintained for each process (or program) by the operating system.

There are three pre-defined file descriptors. You do not need to open these yourself; the operating system provides these automatically:

- File descriptor 0: refers to standard input (usually the terminal keyboard)
- File descriptor 1: refers to standard output (usually the terminal screen)
- File descriptor 2: refers to standard error output (usually the terminal screen)

**Important:** When a process (or program) begins its execution, it starts out with three opened files — i.e. the three pre-defined file descriptors 0, 1, and 2. When the process opens its first file, it will then be attached with file descriptor 3.

### 2.2.2 Opening and creating files

In order to read from or write to a file (including the standard input and output), that file must first be opened. The routine (or function) used to *open* a file or to *create* a file for reading and/or writing is called `open()`.

The `open()` function takes two arguments:

- a character string consists of the path name of a file to be opened;
- a set of flags or integer constants specifying how the file is to be opened;

The man page for `open` (`man 2 open`) illustrates how this function is declared internally inside the system libraries, and what headers you must `#include` to use it.

Function prototype for the `open` system call:

```
1  #include <sys/types.h>
2  #include <sys/stat.h>
3  #include <fcntl.h>
4
5  int open(const char *pathname, int flags);
```

The second argument is constructed by OR-ing together<sup>3</sup> a number of flags or constants (see below) defined in the header file `sys/fcntl.h` for System V systems, and `sys/file.h` for BSD (Berkeley Software Distribution) systems.

A subset of the flags (constants) that control how a file is to be opened are shown in the table on the next page. Note that the first three flags shown (i.e. `O_RDONLY`, `O_WRONLY` and `O_RDWR`) are mutually exclusive.

The `open` routine returns a file descriptor for the file if it is opened successfully. Otherwise, the value of `-1` is returned if it fails to open that file. An error code indicating the reason for failure is stored in the external variable `errno` defined in the header file `errno.h`. You can print out the error message with the `perror` function<sup>4</sup>. Note that `perror` is not a system call.

<sup>3</sup>Use the bitwise-or operator: `|` to combine multiple flag values together if needed.

<sup>4</sup>See `open_perror.c` for an example.

Flag (Constant)	Description
O_RDONLY	The file is opened for read only
O_WRONLY	The file is opened for write only
O_RDWR	The file is opened for both read and write
O_APPEND	The file is opened in append mode
O_CREAT	Create the file if it does not exist; the mode argument must be supplied to specify the access permission on the file
O_EXCL	Check if the file to be created is already exists; if already exists, open() will fail
O_NONBLOCK or O_NDELAY	The file is opened in non-blocking mode
O_TRUNC	If the file is opened for writing, truncate the length of the file to zero

Often times, the opened file should be closed once a process (or program) has finished using that file. The routine to *close* a file is called `close()`, which takes in only one argument representing the file descriptor of the file to be closed. If the file was closed successfully, the value of 0 is returned; otherwise, the value of -1 is returned if an error occurs<sup>5</sup>.

Function prototype for the `close` system call:

```
1  #include <unistd.h>
2
3  int close(int fd);
```

### 2.2.3 Reading and writing files

With the low-level I/O interface, the `read()` and `write()` routines are used for reading from and writing to a file respectively.

Function prototypes for the `read` and `write` system calls:

```
1  #include <unistd.h>
2
3  ssize_t read(int fd, void *buf, size_t count);
4  ssize_t write(int fd, const void *buf, size_t count);
```

---

<sup>5</sup>See `open_perror.c` for an example.

The `read()` system call attempts to read up to `count` bytes from the file referred to by the file descriptor `fd` into the buffer pointed by `buf`. The function returns the number of bytes actually read if no error occurs; otherwise the value of `-1` is returned and the external variable `errno` is set with the error code. If the return value is `0`, that indicates the end of the file has been reached and there is no data left to be read.

The `write()` system call writes up to `count` bytes from the buffer pointed by `buf` to the file referred to by the file descriptor `fd`. If the write is successful, the number of bytes actually written is returned; otherwise the value of `-1` is returned if an error occurs and `errno` is set accordingly. (If nothing was written, the return value is `0`.)

When using these low-level functions, the programmers have to decide on the size of the buffer in bytes (as indicated by the third argument `count` in both functions). Note that each time the `read()` or `write()` function is executed, the operating system will access the disk (or other I/O device). If the buffer size is set to `1`, meaning that you would only be able to read or write one byte (one character) at a time — the overall execution of your program is less efficient (especially when dealing with a large file) compared to reading or writing a whole block of characters at a time. In many cases the difference is not noticeable however.

#### 2.2.4 Using `write()` with strings

While many C library functions use the C language convention of 'null-terminated' strings (using a `'\0'` character to mark the end of the string in memory), the underlying operating system does not assign any special meaning to the null `'\0'` character.

Therefore, when making a `write()` system call to the operating system, you must explicitly specify the number of characters to be written, regardless of whether the string is already null-terminated.<sup>6</sup> The `strlen()` function defined in `<string.h>` can be useful here: this function counts the number of characters in a (null-terminated) string by finding the position of the null character.

**Common mistake:** Beware of trying to use the `sizeof()` function to get the length of a string: depending on how your string is defined, you might be mistakenly getting the size of a `char*` (memory addresses are always 4 bytes in size on a 32-bit machine) or the size of an entire array, rather than the number of meaningful characters defined within it. In general, the `sizeof()` function is not the right tool for this job. It's only really meant for getting the memory size of data types rather than values themselves.

---

<sup>6</sup>In some cases, such as image and video processing, the null character may even represent useful data rather than the end of the string.

The following example shows how to use the `write()` system call to print a string to standard output (which is file descriptor 1)<sup>7</sup>:

```
1 /* write.c */
2
3 char* outputString = "Hello , world!\n";
4 write(1, outputString, strlen(outputString));
5
6 // Can we put the whole thing in one line like with a printf()?
7 write(1, "This also works\n", 16);
8
9 /*...but note that hard-coding the length like this is
10 not considered ideal coding practice for readability! */
```

### 2.2.5 Moving around files

Often times, we need to be able to move around within a file to access the data stored at a specific location in the file. Each file is associated with a value, known as the *file offset*. It is often set to 0 indicating the beginning of the file when a file is opened or created.

The value of the file offset can be obtained or re-set by using `lseek()`.

Function prototype for the `lseek` system call:

```
1 #include <sys/types.h>
2 #include <unistd.h>
3
4 off_t lseek(int fd, off_t offset, int whence);
```

The `lseek()` function re-positions the file offset (`offset`) of a file referred to by the file descriptor `fd`, based on the position in the file specified by the last argument in the function `whence`. The values that can be assigned to `whence` are as follows:

Offset (Constant)	Description
SEEK_SET	The file offset is set to <code>offset</code> bytes from the beginning of the file
SEEK_CUR	The file offset is set to <code>offset</code> bytes from its current position
SEEK_END	The file offset is set to <code>offset</code> bytes from the end of the file

---

<sup>7</sup>See `write.c` for the full source code.

So, if you were to move to the beginning or the end of a file:

```
1  lseek(fd, 0, SEEK_SET);    /* move to the beginning of the file */
2  lseek(fd, 0, SEEK_END);    /* move to the end of the file */
```

To obtain the value of the current offset of the file:

```
1  off_t current;
2  current = lseek(fd, 0, SEEK_CUR);
```

Note that the `lseek()` returns the new offset value in bytes from the beginning of the file if it runs successfully. If there is an error, the value of `-1` is returned. (The `errno` variable is set to indicate the error.)

## **2.3 Pre-class Quiz Checkpoint (2 marks)**

Now that you have completed reading this section, please attempt the pre-class quiz for Laboratory 2 on Moodle before you attend the class.

The pre-class quiz is worth 2 out of 10 raw marks for the in-class assessment of each module. The tasks that follow are worth 8 out of 10 raw marks. In total, this class's assessment tasks are worth 1% of your total final mark for the unit.

## 3 Laboratory Tasks (8 marks)

### 3.1 Task 1 (2 marks)

Given the following program `task1.c`, compile and run it. Describe in detail (3-4 sentences) what the program does, as a comment at the beginning of the code.

```
1  /* task1.c */
2
3  #include <sys/types.h>
4  #include <sys/stat.h>
5  #include <sys/file.h>      /* change to <sys/fcntl.h> for System V */
6  #include <unistd.h>
7  #include <string.h>
8  #include <stdlib.h>
9
10 struct Record {
11     int unitid;
12     char unitcode[8];
13 };
14
15 char *unitcodes[] = {"FIT2100", "FIT1047", "FIT3159", "FIT3142"};
16
17 int main(int argc, char *argv[])
18 {
19     int i, outfile;
20     struct Record eachrec;
21
22     if ((outfile = open("recordfile", O_WRONLY | O_CREAT | O_TRUNC, 0664)) <
23         0) {
24         exit(1);
25     }
26
27     for (i = 3; i >= 0; i -= 2) {
28         eachrec.unitid = i;
29         strcpy(eachrec.unitcode, unitcodes[i]);
30
31         lseek(outfile, (long) i * sizeof(struct Record), SEEK_SET);
32         write(outfile, &eachrec, sizeof(struct Record));
33
34         if (i == 1) {
35             i = 4;
36         }
37     }
38
39     close(outfile);
40     exit(0);
41 }
```



### 3.2 Task 2 (3 marks)

By modifying the program that you have understood in Task 1 (Section 3.1), **write a C program that uses system calls for I/O and does the following:**

- Reads the `Record` structs from the `recordfile` file in the following order: the second record (1), the first record (0), the last or fourth record (3), and the third record (2).
- Prints all unitcodes out on the terminal screen
- Writes the unitcodes with an *odd* unitid to a file called `odd`.
- Writes the unitcodes with an *even* unitid to a file called `even`.

### 3.3 Task 3 (3 marks)

Download `task3.c` from Moodle. This `task3.c` file is a partially completed C program, that takes two file names — `file1` and `file2` — as its command-line arguments. It opens the first file (`file1`) for writing and the second file (`file2`) for reading. The contents of the second file (`file2`) is appended at the end of the first file (`file1`).

**Complete the missing code segment in the program as indicated by the comment — `/* CODE HERE */`.** To test the program, create two files and name them as `file1` and `file2`. The contents of each file is given below by using the `less` command in Unix.

```
1  $ less file1
2  Line 1: the first sentence in file1
3  Line 2: the second sentence in file1
4  Line 3: the third sentence in file1
5
6  $ less file2
7  Line 1: the first sentence in file2
8  Line 2: the second sentence in file2
9  Line 3: the third sentence in file2
```

Note: To run or test the program in Unix, use the following command:

```
1  $ ./appendfile file1 file2
```

### 3.4 Task 4 (Extension: 0 marks)

Try running NetData (found on your desktop under the Dev Tools directory) while your program is running. Locate the graph for disk I/O. Do you see any disk activity (in or out) when running your Task 3 program? Our files are quite small. How does it compare to the amount of disk I/O activity you see when loading a web browser or other large application?

## 4 Postscript: Redirects and Pipes

You might be wondering why Unix provides two different output streams (standard output, and also standard error). The answer is that Unix systems generally provide a lot of flexibility in how input and output are managed. The following operators are provided by the bash command shell, and work with any program, even if your program is not written to use them:

```
1 $ ./some_program > outputfile.txt
2 [Standard output produced by the program is redirected into a file instead]
3
4 $ ./some_program 2> errorfile.txt
5 [the program's error messages are redirected into a file instead of the
   terminal]
6
7 $ ./some_program < input_commands.txt
8 [the program takes its input from a file, not from the keyboard]
9
10 $ ./some_program | ./another_program
11 [the standard output from some_program is used as input for another_program
   . This is called a 'pipe'.]
```

These operators are invisible to the running program. The bash command shell does not pass them to your program as arguments.

Another example involves using the special `/dev/null` device file. This is a special virtual device on a Linux system. Anything written to it is simply discarded. Suppose you want to hide error messages when running a certain program, and only want to see the meaningful output:

```
1 $ ./noisy_program 2> /dev/null
```